

Non-Linear Regression

Václav Šmíd

April 24, 2018

Linear regression and OLS

Fit by a linear function:

$$\begin{array}{cccc} y_1 & = & ax_1 & + b_1, & + e_1 \\ y_2 & = & ax_2 & + b_1 & + e_2, \\ \vdots & & \vdots & & \vdots \end{array}$$


Minimize

$$\sum_i e_i^2 = \sum_i (y_i - ax_i - b)^2$$

Solution:

$$\frac{d(\sum_i e_i^2)}{d\theta} = 0.$$

$$\hat{\theta} = (X^T X)^{-1} X^T \mathbf{y}.$$




Limit of fixed bases

Fit by a linear function:

$$y_i = a\phi_1(x_i) + b\phi_2(x_i) + e_i$$

What are the basis function?



Limit of fixed bases

Fit by a linear function:

$$y_i = a\phi_1(x_i) + b\phi_2(x_i) + e_i$$

What are the basis function?

Basis functions are functions of data:

$$y_i = a\phi_1(\psi_1, x_i) + b\phi_2(\psi_2, x_i) + e_i$$


where ϕ_j are non-linear functions.

Estimating new set of parameters

$$\theta = [a, b, \psi_1, \psi_2]$$

$$\frac{d(y_i - a\phi_1(\psi_1, x_i) + b\phi_2(\psi_2, x_i))}{d\theta} = 0.$$

$$\hat{\theta} = ?$$



Neural networks


Feed forward NN:

$$z = \sigma_1 (W_1 x + b_1),$$

$$y = \sigma_2 (w_m z_m + b_m)$$

with vector-valued

- **activation** functions $\sigma_j()$,
- **weights** w_j
- **biases** b_j .



For Gaussian noise, MSE (mean square error) loss function:

$$L = \sum_{i=1}^n (y_i - \sigma_1 (w_1 \sigma_2 (\cdots) + b_1))^2.$$

finding $\theta = [w_1, b_1, w_2, b_2, \dots]$ by

Neural networks


Feed forward NN:

$$z = \sigma_1 (W_1 x + b_1),$$

$$y = \sigma_2 (w_m z_m + b_m)$$

with vector-valued

- **activation** functions $\sigma_j()$,
- **weights** w_j
- **biases** b_j .



For Gaussian noise, MSE (mean square error) loss function:

$$L = \sum_{i=1}^n (y_i - \sigma_1 (w_1 \sigma_2 (\dots) + b_1))^2.$$

finding $\theta = [w_1, b_1, w_2, b_2, \dots]$ by gradient descent method

$$\hat{\theta}^{(\tau+1)} = \hat{\theta}^{(\tau)} - \eta \nabla L(\hat{\theta}^{(\tau)}),$$

Regularization where η is the (small) learning rate.

Example



Trivial NN with one hidden layer:

$$y_i = \sum_{i=1}^6 w_{2,i} \tanh(w_{1,j}x_i + b_{1,j}) + b_2,$$


tanh activation function on hidden layer and linear activation function on output.

Training:

1. random initialization,
2. 50000 steps,
3. rate $\eta = 0.001$,



Convergence issues



- ▶ Step-size tuning, schedule, annealing, ...
- ▶ Better gradient (ADAM) [Kingma, Ba, 2014].
- ▶ Higher order methods (half-quadratic approximation, Hessian, LMBFS)

Preventing overfitting

1. ridge regression

$$L = \sum_{i=1}^m (y_i - \sigma_1(w_1 \sigma_2(\dots) + b_1))^2 + \alpha \sum_{k,l} w_{k,l}^2.$$

2. Automatic relevance determination (Laplace)

$$L = \sum_{i=1}^m (y_i - \sigma_1(w_1 \sigma_2(\dots) + b_1))^2 + \sum_{k,l} \alpha_{k,l} w_{k,l}^2.$$

3. Stochastic gradient descent

$$\nabla L(\hat{\theta}) = \sum_{i=1}^n \nabla L(\hat{\theta}, x_i, y_i) \approx \nabla L(\hat{\theta}) = \sum_{i=\mathcal{I}} \nabla L(\hat{\theta}, x_i, y_i),$$

where \mathcal{I} is a random subsample of $\{1, \dots, n\}$.

Points

	points
Ridge regression	10
ADAM (own implementation)	10
Stochastic Gradient Descent	10
Relevance determination	30